



ilifu Online Training

Dr Jordan Collier

ilifu Support Astronomer, IDIA, Department of Astronomy, University of Cape Town
Adjunct Fellow, Western Sydney University



UNIVERSITY OF CAPE TOWN
IYUNIVESITHI YASEKAPA • UNIVERSITEIT VAN KAAPSTAD

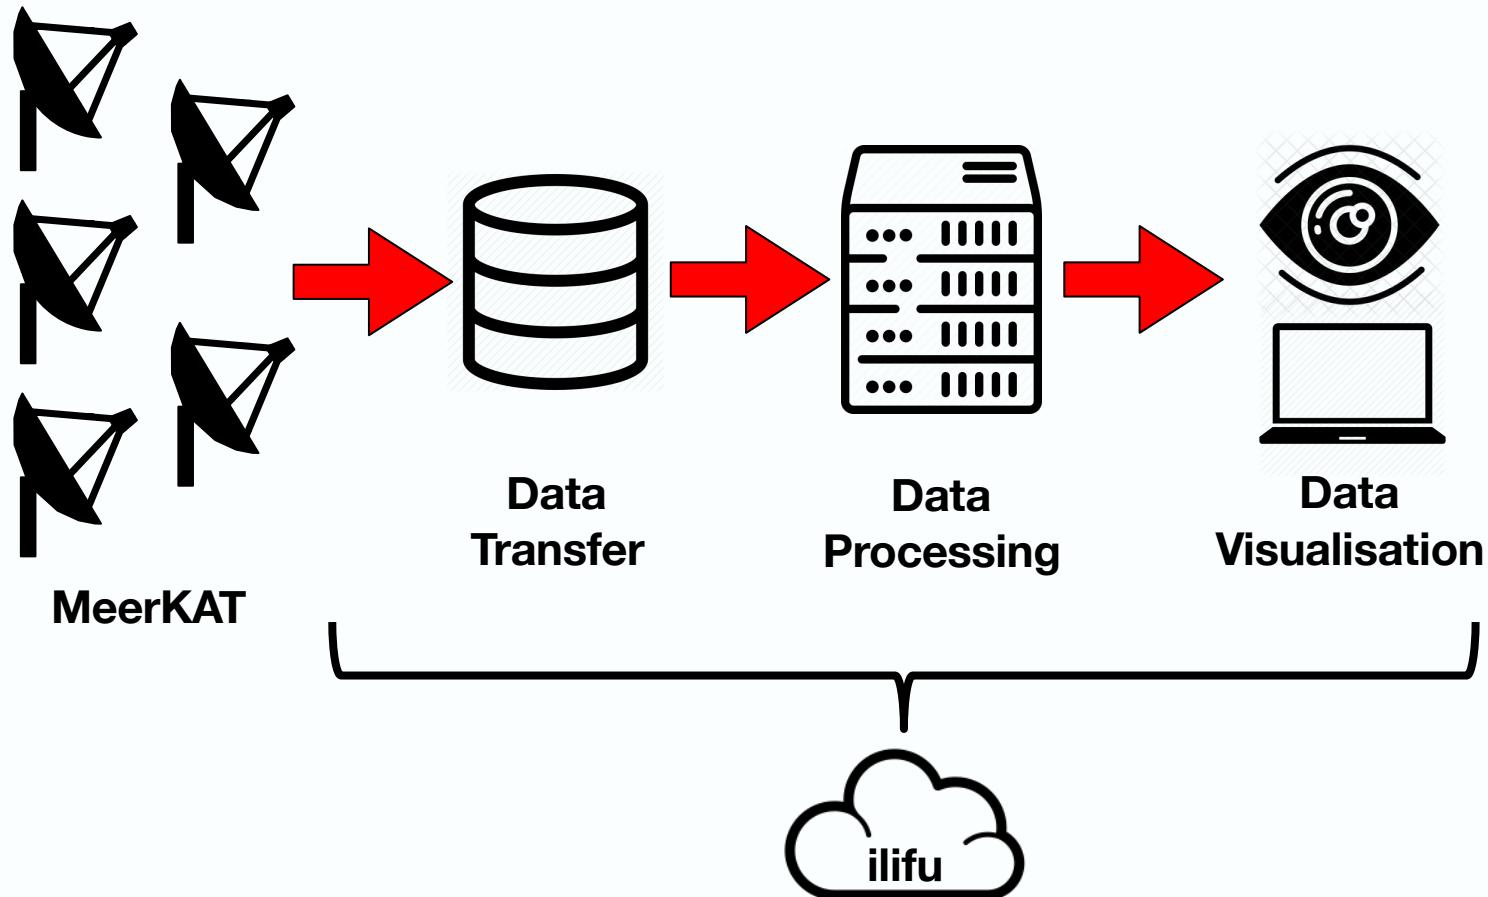


IDIA Inter-University Institute
for Data Intensive Astronomy

WESTERN SYDNEY
UNIVERSITY



ilifu (Biased) Dataflow



Data Transfer

- http://docs.ilifu.ac.za/#/data/data_transfer
- Do not use head node!
- Don't write large files to /users, only scripts/config files
- transfer.ilifu.ac.za
 - For scp and rsync
 - e.g. `$ scp /path/to/file/<filename> <username>@transfer.ilifu.ac.za:/idia/users/<username>/scripts/`
- Globus
 - Faster than scp and rsync

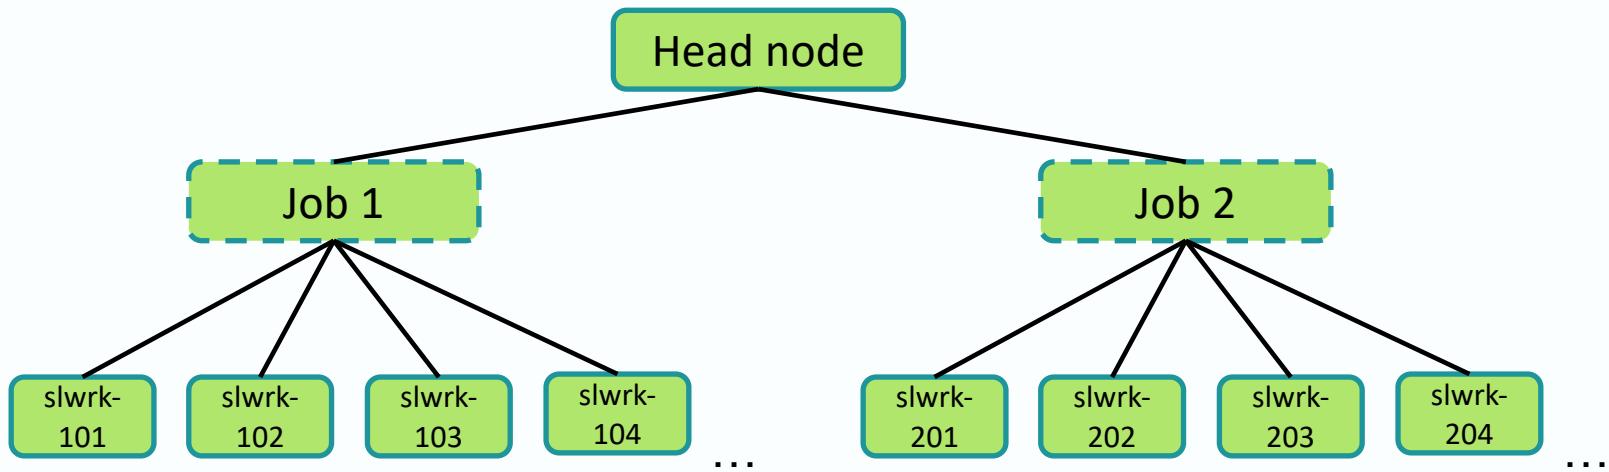
Data Transfer: Globus

- GridFTP transfer service setup between SARAQ and IDIA
 - Currently extending to collaborators in UK, Netherlands, Italy, India, etc
- User-friendly globus connect built on top of gridFTP
 - GUI/web app or CLI
- Can connect any arbitrary end points
 - Server (DTN), desktop, etc for Mac OS X, Windows & Linux
 - Offers user-friendly but computationally efficient transfer service
- Demo time!



Parallelism

- Oxford definition for parallel processing
 - *a mode of operation in which a process is split into parts, which are executed simultaneously on different processors attached to the same computer.*
- The work is partitioned into smaller jobs, sometimes with a partition of the dataset



Parallelism

- Can be achieved on a single machine / node
 - Distributes work over many CPUs
 - Typically implemented using OpenMP
- Or over multiple machines / nodes
 - Distributes work over many tasks, over 1+ nodes
 - Each given amount of memory to use
 - Generally requires a cluster
 - Typically implemented using OpenMPI
 - Requires a message passing interface (MPI) wrapper
 - mpirun, aprun, srun (SLURM), mpicasa (CASA)
- Managed on ilifu by SLURM



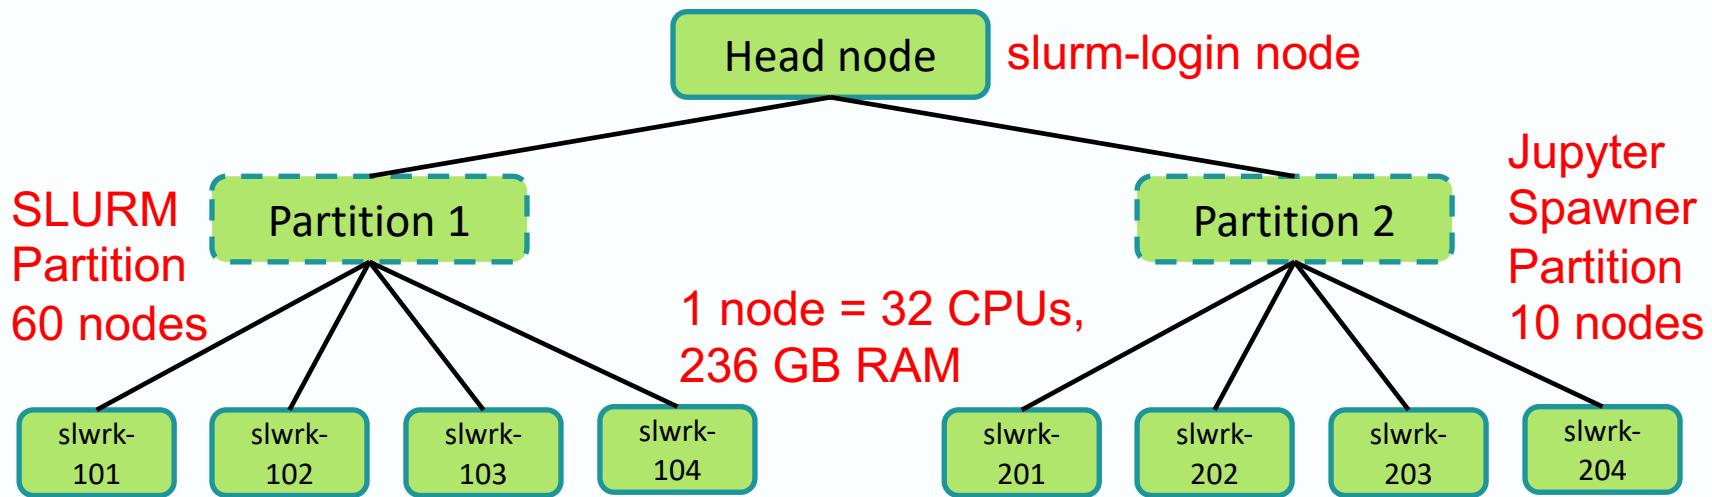
Parallelism

- Implementing a normal job in SLURM
 - Will only use 1 CPU, 1 task, and 1 node
 - Default for many processes
- Implementing an OpenMP job in SLURM
 - Need to use >1 CPU, while nodes & tasks must be 1 (unless also using MPI)
 - cpus-per-task
 - May need to export OMP_NUM_THREADS
- Implementing an MPI job in SLURM
 - Need to use >1 task, while nodes and CPUs can be 1
 - nodes
 - ntasks-per-node
 - cpus-per-task
- See upcoming examples



SLURM

- http://docs.ilifu.ac.za/#/getting_started/submit_job_slurm
- Head node (job submission & management)
- Partitions
- Compute nodes



SLURM

- http://docs.ilifu.ac.za/#/getting_started/submit_job_slurm
- Head node (job submission & management)
- Partitions
- Compute nodes

ssh - shell terminal

```
* Support: https://ubuntu.com/advantage
System information as of Fri Aug 23 11:36:57 SAST 2019
System load: 0.49      Users logged in: 8
Usage of /: 35.9% of 21.15GB  IP address for ens3: 192.168.100.39
Memory usage: 5%       IP address for ens4: 10.102.26.97
Swap usage: 0%         IP address for ens5: 10.102.28.133
Processes: 396

* Keen to learn Istio? It's included in the single-package MicroK8s.

  https://snapcraft.io/microk8s

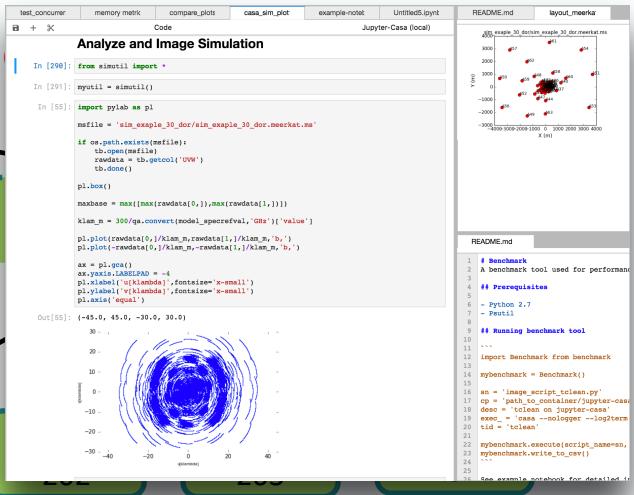
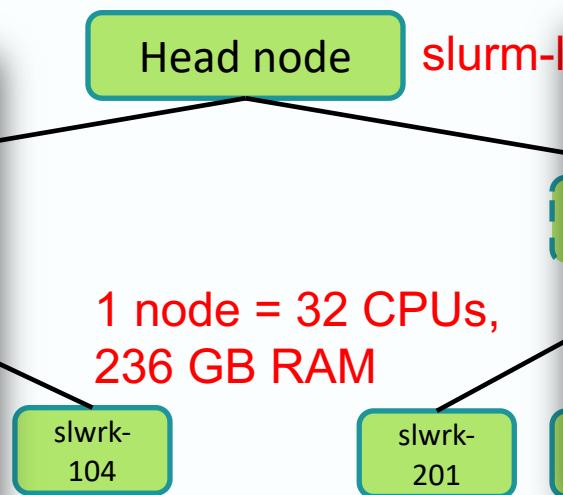
Get cloud support with Ubuntu Advantage Cloud Guest:
  http://www.ubuntu.com/business/services/cloud

* Canonical Livepatch is available for installation.
- Reduce system reboots and improve kernel security. Activate at:
  https://ubuntu.com/livepatch

170 packages can be updated.
75 updates are security updates.

Last login: Fri Aug 23 09:00:21 2019 from 196.11.235.232
jeremy@slurm-login:~$ sinfo
PARTITION          AVAIL TIMELIMIT NODES STATE NODELIST
Main*             up 14:00:00:9      8  mix slwrk-[106-113]
Main*             up 14:00:00:9      14 alloc slwrk-[101,104-105,114-124]
Main*             up 14:00:00:9      38  idle slwrk-[102-103,125-160]
JupyterSpawnerONLY up infinite      4  mix slwrk-[201-202,205,209]
JupyterSpawnerONLY up infinite      4  alloc slwrk-[206-208,210]
JupyterSpawnerONLY up infinite      2  idle slwrk-[203-204]
jeremy@slurm-login:~$ sbatch compute_job.sh
```

\$ ssh <username>@slurm.ilifu.ac.za



SLURM – user commands

- <https://slurm.schedmd.com/>

```
$ squeue #shows all jobs in slurm queue
```

```
$ squeue -u <username> #shows your jobs
```

```
$ sinfo #shows partitions and resources
```

```
$ sbatch slurm_job_script.sh #submit job to  
slurm queue
```

```
$ sbatch --help #describes input parameters
```

```
$ scancel <jobid> #cancels job
```

```
$ scontrol show jobid <jobid> -dd #shows  
useful information about your job
```

SLURM - running a job

```
$ cat slurm_job_script.sh  
#!/bin/bash  
  
#SBATCH --job-name='demo_job'  
#SBATCH --ntasks=1  
#SBATCH --cpus-per-task=4  
#SBATCH --mem-per-cpu=4GB  
#SBATCH --output=demo_job-%j-stdout.log  
#SBATCH --error=demo_job-%j-stderr.log  
  
echo "Submitting demo SLURM job"  
  
singularity exec /idia/software/containers/SF-PY3-bionic.simg python myscript.py
```



Describe job parameters/resources

container

software

script

```
$ sbatch slurm_job_script.sh #submit job to slurm queue
```

what's being executed with above parameters

SLURM - running a job

```
$ cat slurm_job_script.sh  
#!/bin/bash  
  
#SBATCH --job-name='demo_job'  
#SBATCH --ntasks=1  
#SBATCH --cpus-per-task=4  
#SBATCH --mem-per-cpu=4GB  
#SBATCH --output=demo_job-%j-stdout.log  
#SBATCH --error=demo_job-%j-stderr.log  
  
echo "Submitting demo SLURM job"  
singularity exec /idia/software/containers/casa-stable.img casa -c myscript.py
```



Describe job parameters/resources

container

software

script

```
$ sbatch slurm_job_script.sh #submit job to slurm queue
```

what's being executed with above parameters

SLURM - running an interactive job

- http://docs.ilifu.ac.za/#/getting_started/submit_job_slurm?id=interactive-slurm-session

```
$ salloc #opens bash shell session on compute node,  
#default 1 task, 1 CPU, 8GB RAM
```

```
$ salloc --cpus-per-task=4 --mem=64GB --pty singularity  
exec /idia/software/containers/casa-stable.simg casa  
#opens interactive CASA session on compute node,  
#with 1 task, 4 CPUs, 64GB RAM,  
#using Singularity CASA container
```

```
$ salloc --help  
#view salloc help docs for input parameters
```

SLURM - running an interactive job

- With X11 support for viewing GUI

```
$ ssh-add -K / -k #this may need to be run before ssh, if you've  
copied keys from an old machine (-K for Mac OS X, -k for Linux)
```

```
$ ssh -YA <username>@slurm.ilifu.ac.za
```

#important to include the -YA parameters when
logging into the SLURM login node

```
$ salloc --qos qos-interactive
```

#opens bash shell session on
compute node with X11 support
#default 1 task, 1 CPU, 8GB RAM

```
$ salloc --cpus-per-task=4 --mem=64GB --qos qos-interactive
```

SLURM – general advice

- Before running a large job, identify the available resources
 - Use sinfo
- Don't hog the cluster
 - Reduce your allocation if necessary / possible
- For interactive jobs, use salloc, or srun --pty bash
- Don't ssh onto the worker nodes!
 - This takes away from other peoples resources, even when they've reserved those resources with SLURM
 - We will kill your job if you ssh directly onto the node!
- Demo time!

IDIA MeerKAT Pipeline – A Good Framework

- Parallelised package for HPC processing (SLURM + cluster)
 - Uses multi-measurement sets (MMS) to parallelise across a cluster
- Outputs calibrated data (push of button!)
- HPC-friendly – dynamically uses resources & submits to queue
- Each job/script is a logical step that does / doesn't use MPI, and optionally uses a different container
- User can insert your scripts at start, middle or end
- <https://idia-pipelines.github.io/docs/processMeerKAT>
- Demo time!

CARTA

- Cube Analysis and Rendering Tool for Astronomy
 - IDIA (South Africa) – NRAO (US) – ASIAA (Taiwan)
 - Cloud-based Visual analytic of remote large image cubes
 - Supports many image formats: FITS, CASA, Miriad, and HDF5
- CARTA on ilifu
 - http://docs.ilifu.ac.za/#/astronomy/astronomy_software?id=carta
 - <https://carta.idia.ac.za> – login with same credentials as Jupyter
 - The only visible files are in /carta_share/
 - User should write IDIA HDF5 file to this filesystem
 - Even for continuum (single channel) image
 - `srun /carta_share/hdf_convert/run_hdf_converter image.fits /carta_share/users/<username>/image`
- Demo time!



THANK YOU

Dr Jordan Collier

ilifu Support Astronomer, IDIA
Department of Astronomy,
University of Cape Town

Adjunct Fellow, Western Sydney University
School of Computing, Engineering and
Mathematics

Jordan@idia.ac.za
+27 664 343 953 (RSA Mobile)
+61 414 443 622 (AU Mobile / WhatsApp)



Inter-University Institute
for Data Intensive Astronomy



UNIVERSITY OF CAPE TOWN
IYUNIVESITHI YASEKAPA • UNIVERSITEIT VAN KAAPSTAD

