A photograph of a radio telescope array at night, with several large parabolic dishes illuminated against a starry sky. The dishes are mounted on tall, white, cylindrical pedestals. The foreground shows a dark, flat landscape with some low-lying vegetation and a yellow-painted curb.

ilifu Online Training – Advanced slurm

Tinus Cloete

System Administrator & User Support, ilifu
University of Western Cape, Sept 2025



- Interactive Jobs
 - Persistent Terminals (`tmux`)
- Advanced Slurm monitoring:
 - How busy is the cluster? (`sinfo`)
 - Where are my jobs in the queue? (`squeue`)
 - Which jobs have I submitted in the past and what happened? (`sacct`)
- Advanced Slurm job submission:
 - Slurm Arrays
 - Slurm Job Dependencies

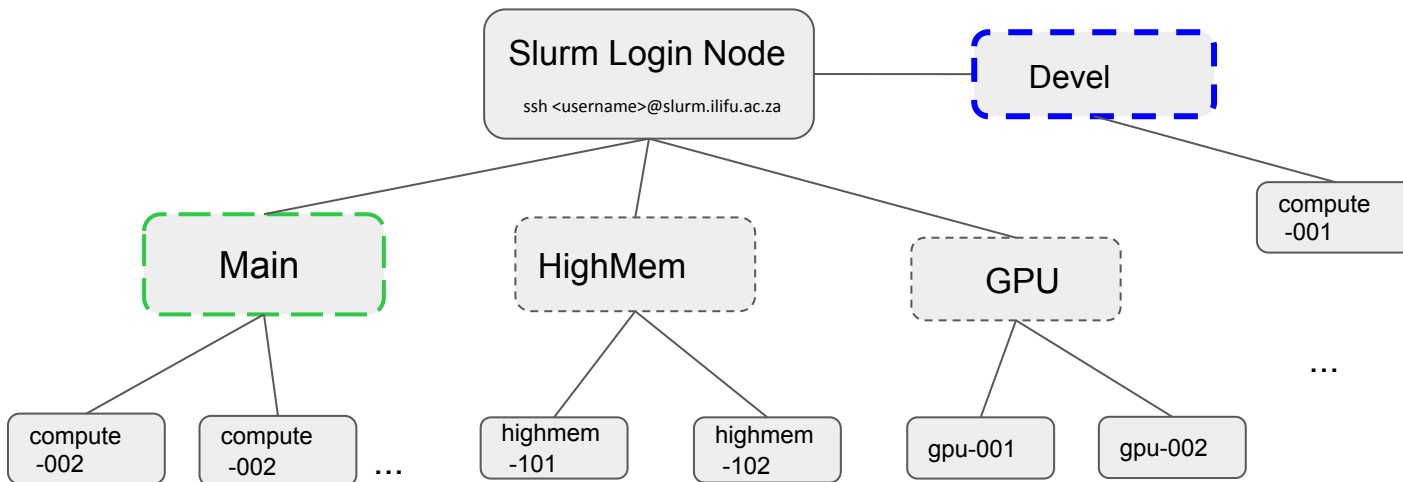
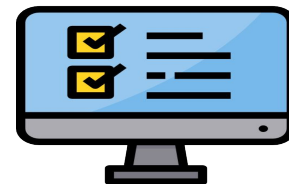


SLURM

http://docs.ilifu.ac.za/#/tech_docs/running_jobs?id=specifying-resources-when-running-jobs-on-slurm



Main	HighMem	GPU	Devel
~85 nodes	8 nodes	7 nodes	1 node
32 CPUs / 232 GB RAM	32 CPUs / 480 GB RAM	2 GPUs / 232 GB RAM	32 CPUs / 237 GB RAM



SLURM - running an interactive job



Login to Slurm Login Node

```
$ ssh <username>@slurm.ilifu.ac.za
```

```
$ sinteractive
```

- Devel partition

```
tcloete@slurm-login:~$ sinteractive
Starting interactive Slurm session.
srun: job 9387238 queued and waiting for resources
srun: job 9387238 has been allocated resources
tcloete@compute-001:~$ |
```



The sinteractive script is a wrapper script we created for calling Slurm's *srun* command with useful default parameters.

```
$ srun --pty bash
```

- Main partition

```
srun -p Main --pty bash
```

- Main partition



UNIVERSITY of the
WESTERN CAPE

SLURM - running an interactive job



Login to Slurm Login Node

```
$ ssh <username>@slurm.ilifu.ac.za
```

Defaults: 1 CPU for 3 hours:

```
$ sinteractive -c 1 -time 03:00
```

Example: 5 CPU for 5 days (maximum):

```
$ sinteractive -c 5 -time 5-00:00
```

Note: All resources on the Devel partition are shared. Including CPU and memory.

For interactive jobs requiring dedicated resources, please use: **srun** on the Main partition



SLURM - running an interactive job (x11 support)



Login to Slurm Login Node

```
$ ssh -Y <username>@slurm.ilifu.ac.za
```

```
$ sinteractive --x11
```

Testing X11 Support:

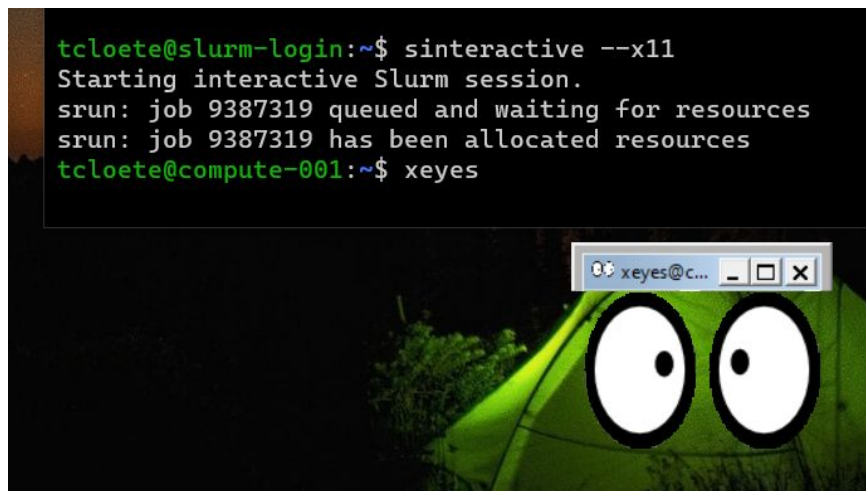
```
$ xeyes
```

OR

```
$ xmessage 'hello'
```

Windows users:

Install X11 server:
e.g. Xming or Vcxsrv



tmux — terminal multiplexer

- Persistent Connections
- Runs Multiple Terminals

Note: Persistent terminals can be lost if the Slurm-login node had to be restarted

Commands

- | | |
|----------------|-----------------------------|
| \$ tmux | # start a new tmux session |
| \$ tmux attach | # Attach to running session |
| | |
| \$ tmux ls | # List active sessions |
| \$ tmux --help | # show help |

```
tcloete@slurm-login:~$ echo "terminal 1"
terminal 1
tcloete@slurm-login:~$

tcloete@slurm-login:~$ echo "terminal 2"
terminal 2
tcloete@slurm-login:~$ |

tcloete@slurm-login:~$ echo "terminal 3"
terminal 3
tcloete@slurm-login:~$

tcloete@slurm-login:~$ echo "terminal 4"
terminal 4
tcloete@slurm-login:~$

[10] 0: bash* "tcloete@slurm-login: " 08:58 20-Mar-24
```



Keyboard shortcuts

Ctrl+b d # detach session (can reconnect)

Ctrl+b x # kill terminal pane

Ctrl+b % # split the screen horizontally

Ctrl+b " # split the screen vertically

Ctrl+b <arrow key> # switch to the pane in arrow direction

Ilifu docs Reference:

https://docs.ilifu.ac.za/#/tech_docs/running_jobs?id=persistent-terminals

Video Tutorial (Jeremy Howard): <https://youtu.be/0pWjZByJ3Lk?t=2474>



DEMO TIME!



- Interactive Jobs
- Persistent Terminals (`tmux`)
- Advanced Slurm monitoring:
 - How busy is the cluster? (`sinfo`)
 - Where are my jobs in the queue? (`squeue`)
 - Which jobs have I submitted in the past and what happened? (`sacct`)
- Advanced Slurm job submission:
 - Slurm Arrays
 - Slurm Job Dependencies



Slurm Commands	Before Running a Job	Starting a New Job	Monitoring a Running Job	After Jobs have Completed
sinfo	X			
squeue	X		X	
sbatch		X		
srun		X		
scontrol*			X	
sacct				X

* Shows all information about a pending/running job:

\$ scontrol show job <jobID>



sinfo: How busy is the cluster?



```
$ sinfo
```

```
PARTITION AVAIL TIMELIMIT NODES STATE
Main*      up 14-00:00:0      1 drain*
Main*      up 14-00:00:0     19 mix
Main*      up 14-00:00:0     17 alloc
Main*      up 14-00:00:0     48 idle
```

List partitions and their specs

```
$ sinfo -0 "partition,available,cpus,nodes,memory,statecompact"
```

```
PARTITION AVAIL CPUS NODES MEMORY STATE
Main*      up    32     1   237568 drain*
Main*      up   32+    19  237568+ mix
Main*      up    32    17  237568+ alloc
Main*      up   32+    48  237568+ idle
Jupyter    up    32     1   237568 drng
Jupyter    up    32     6   237568 mix
Jupyter    up    32     2   237568 alloc
Jupyter    up    32     1   237568 idle
JupyterGPU up    32     1   237568 mix
JupyterGPU up    32     1   237568 alloc
HighMem    up    96     1  1544192 mix
HighMem    up    32     2   515604 idle
GPU         up    24     1   237568 down*
GPU         up    32     3   237568 mix
GPU         up    32     1   237568 alloc
GPU         up    48     2   362582 idle
GPUV100    up    24     1   237568 down*
Devel      up    32     1   237568 alloc
```

For a list of sinfo columns:

```
$ man sinfo
```



UNIVERSITY of the
WESTERN CAPE

queue: Where is my job in the queue?



```
$ squeue -u $USER
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
9402788	Devel	bash	tcloete	R	0:32	1	compute-001
9402721	Main	bash	tcloete	PD	0:00	1	(Resources)
9402719	Main	bash	tcloete	R	7:07	1	compute-260

```
$ squeue -u $USER -o JobID,Partition,State,NumCPUs,MinMemory,NodeList,ReasonList
```

JOBID	PARTITION	STATE	CPUS	MIN_MEMORY	NODELIST(REASON)
9402788	Devel	RUNNING	1	0	compute-001
9402721	Main	PENDING	30	3096M	(Resources)
9402719	Main	RUNNING	30	3096M	compute-260

```
$ squeue -u $USER --start
```

JOBID	PARTITION	NAME	USER	ST	START_TIME	NODES	SCHEDNODES	NODELIST(REASON)
9402721	Main	bash	tcloete	PD	2024-03-25T11:49:02	1	compute-260	(Resources)

sacct: Which jobs have been submitted in the past?



Show jobs submitted by the user in the past 1 day:

```
$ sacct -u $USER
```

Show one line per job (combines job steps):

```
$ sacct -u $USER -X
```

Outputs a list of available column names:

```
$ sacct -e
```

```
$ sacct --helpformat
```

Filter on jobs submitted between a start date and an end date

```
$ sacct -u $USER -X -S 2024-03-25 -E 2025-03-26
```

Show specific columns:

```
$ sacct -u $USER -X -S 2024-03-25 -E 2025-03-26
```

```
-o Jobid,JobName,NCPU,NNodes,Start,End,Elapsed,TimeLimit,State
```

Nicer formatting of columns e.g. %-15 is a column with width of 15 and a left justification

```
$ sacct -u $USER -X
```

```
-o Jobid,JobName%-15,NCPU,NNodes,Start,End,Elapsed,TimeLimit,State%20
```

```
tcloete@compute-001:~/demo/interactive_script$ sacct -u $USER
```

JobID	JobName	Partition	Account	AllocCPUS	State	ExitCode
10042739	bash	Devel	b05-pipel+	1	RUNNING	0:0
10042739.ex+	extern		b05-pipel+	1	RUNNING	0:0
10042739.0	bash		b05-pipel+	1	RUNNING	0:0
10043369	bash	Devel	b05-pipel+	1	COMPLETED	0:0
10043369.ex+	extern		b05-pipel+	1	COMPLETED	0:0
10043369.0	bash		b05-pipel+	1	COMPLETED	0:0

sacct: Which jobs have been submitted in the past?



Example for showing how long jobs ran to help set the Wall time (TimeLimit)

```
$ sacct -u $USER -X  
-o Jobid,JobName%-15,NCPU,NNodes,Start,End,Elapsed,TimeLimit,State%20
```

JobID	JobName	NCPU	NNodes	Start	End	Elapsed	TimeLimit	State
9402635	transfer_MS_15+	0	1	None	2024-03-25T04:39:59	00:00:00	4-00:00:00	CANCELLED by 10028
9402636	partition	23	1	2024-03-25T04:45:05	2024-03-25T04:45:24	00:00:19	00:30:00	FAILED
9402682	partition	23	1	2024-03-25T06:53:06	2024-03-25T07:02:18	00:09:12	00:30:00	COMPLETED
9402683	validate_input	1	1	2024-03-25T07:02:18	2024-03-25T07:02:26	00:00:08	00:30:00	COMPLETED
9402684	flag_round_1	16	1	2024-03-25T07:02:26	2024-03-25T07:17:58	00:15:32	00:30:00	OUT_OF_MEMORY
9403040	partition	23	1	2024-03-25T08:13:18	2024-03-25T08:23:19	00:10:01	00:30:00	COMPLETED
9403041	validate_input	1	1	2024-03-25T08:23:20	2024-03-25T08:23:44	00:00:24	00:30:00	COMPLETED
9403042	flag_round_1	32	1	2024-03-25T08:23:44	2024-03-25T08:52:29	00:28:45	00:30:00	COMPLETED
9403043	setjy	23	1	2024-03-25T08:52:29	2024-03-25T09:20:13	00:27:44	00:30:00	COMPLETED
9403044	xx_yy_solve	1	1	2024-03-25T09:20:13	2024-03-25T09:42:30	00:22:17	00:30:00	COMPLETED
9403045	xx_yy_apply	23	1	2024-03-25T09:42:30	2024-03-25T09:55:40	00:13:10	00:30:00	COMPLETED
9403046	flag_round_2	32	1	2024-03-25T09:55:40	2024-03-25T10:25:45	00:30:05	00:30:00	TIMEOUT

sacct: Which jobs have been submitted in the past?



Example for showing how much resources jobs used

e.g. MaxRSS = Maximum RAM usage in bytes

Note: Only on job-step level

```
$ sacct -u $USER -S 2024-03-25 -E 2025-03-25 -o  
Jobid%-15,JobName%-15,NCPU,NNodes,Start,MaxRSS,MaxDiskRead,MaxDiskWrite,State%20
```

JobID	JobName	NCPUS	NNodes	Start	MaxRSS	MaxDiskRead	MaxDiskWrite	State
9403045	xx_yy_apply	23	1	2024-03-25T09:42:30				COMPLETED
9403045.batch	batch	23	1	2024-03-25T09:42:30	171434320K	799330.48M	458872.78M	COMPLETED
9403045.extern	extern	23	1	2024-03-25T09:42:30	172K	0.01M	0.00M	COMPLETED
9403046	flag_round_2	32	1	2024-03-25T09:55:40				TIMEOUT
9403046.batch	batch	32	1	2024-03-25T09:55:40	452567280K	276836.01M	15385.23M	CANCELLED
9403046.extern	extern	32	1	2024-03-25T09:55:40	200K	0.01M	0.00M	COMPLETED

- After / during running jobs :
 - jobId is given from sbatch output / squeue

Shows info about job running including working directory

```
$ scontrol show job <jobID>
```

Shows queue with start time (%S)


```
$ squeue --start -u $USER
```

Shows info for multi-CPU jobs

```
$ sacct -o JobID%-15,JobName%-15,Partition,Account,  
Elapsed,NNodes%6,NTASK%6,NCPUS%5,MaxDiskRead,MaxDiskWrite,  
NodeList%20,MaxRSS,CPUTime,State,ExitCode
```

Shows jobs started and completed between these dates

```
$ sacct -S 2021-09-01-09:00 -E 2021-09-14-10:00 -X -o  
Jobid,JobName,Start,End,State
```



- Email parameters

```
$ srun --mail-user=<address> --mail-type=<event_types>  
- Events : BEGIN,END,FAIL,TIME_LIMIT_80
```

- Exclude nodes
 - e.g. problematic nodes (report to ilifu support)

```
$ srun --exclude=compute-[101,101-105]
```



- Specify lower wall-time (default 3 hours) and less memory (default ~7GB) increases chance of job launching immediately


- In steps:

```
$ srun --pty --time=10 --mem=1GB bash
$ singularity shell /idia/software/containers/python-3.6.img
$ python3 job_script.py
```

- In single call:

```
$ srun --pty --time=10 --mem=1GB singularity exec
/idia/software/containers/python-3.6.img python3 job_script.py
```

- Must manually process after this




Topics

- Interactive Jobs
- Persistent Terminals (`tmux`)
- Advanced Slurm monitoring:
 - How busy is the cluster? (`sinfo`)
 - Where are my jobs in the queue? (`squeue`)
 - Which jobs have I submitted in the past and what happened? (`sacct`)
- Advanced Slurm job submission:
 - Slurm Arrays
 - Slurm Job Dependencies



- Array jobs allow quick submission of many similar jobs, each the with same resources, without any manual launch
 - Passes array task ID into script, which changes behavior of each job each time i.e different inputs
 - Can be used to run many related steps simultaneously in a serial process
- Example job array Running, 20 jobs, with 5 run concurrently.



```
#!/bin/bash
#SBATCH --array=1-20%5
#SBATCH --job-name=myarrayjob
#SBATCH --output=logs/%x-%A_%a.out
#SBATCH --error=logs/%x-%A_%a.err
```

```
module load python/3.11.0
```

```
python myscript.py --input $SLURM_ARRAY_TASK_ID
```



Parameter	Substitution / filename pattern	Environment Variables
jobID of running job	%j	SLURM_JOB_ID
Job name	%x	SLURM_JOB_NAME
Job array's master job allocation number	%A	SLURM_ARRAY_JOB_ID
Job array task ID (index) number	%a	SLURM_ARRAY_TASK_ID

For example:

```
#SBATCH --array=1-20%5
```

```
#SBATCH --job-name=myarrayjob
```

```
#SBATCH --output=logs/%x-%A_%a.out
```

With submitted first job id: 1000

%j = 1000,1001...,1019

%x = myarrayjob

%A = 1000

%a = 1,2,3,...,19,20 (order not guaranteed)



- Allow jobs to be scheduled for running, based on the status of a previous job
 - e.g only begin a particular job once previous one successfully completes

Submit `another_job.sh` to SLURM queue, to begin after jobID 1234 successfully completes , or cancel the job if jobID 1234 fails

```
$ sbatch -d afterok:1234 --kill-on-invalid-dep=yes another_job.sh
```

Submit `another_job.sh` to SLURM queue, to begin after jobID 1234 & 5678 completes

```
$ sbatch -d afterany:1234:5678 another_job.sh
```



DEMO TIME!



Do's :

- Run jobs using sbatch rather than interactive jobs
- Identify job resources requirements:
 - No. of nodes and CPUs, amount of RAM and wall-time.
- Remove files that aren't needed
 - /scratch3 folder after data processing is complete
 - Old raw data, temporary products , etc.
- Use Singularity (cannot install software on nodes)
- Use **username@transfer.ilifu.ac.za** for data transfers

Don't:

- Don't run software/heavy processes on login node
- Don't place large files in your home directory (/users)
- Don't transfer using scp/rsync on the login node



Thank you!

Remember our support channels!

support@ilifu.ac.za
<https://docs.ilifu.ac.za>

