

# ilifu Online Training

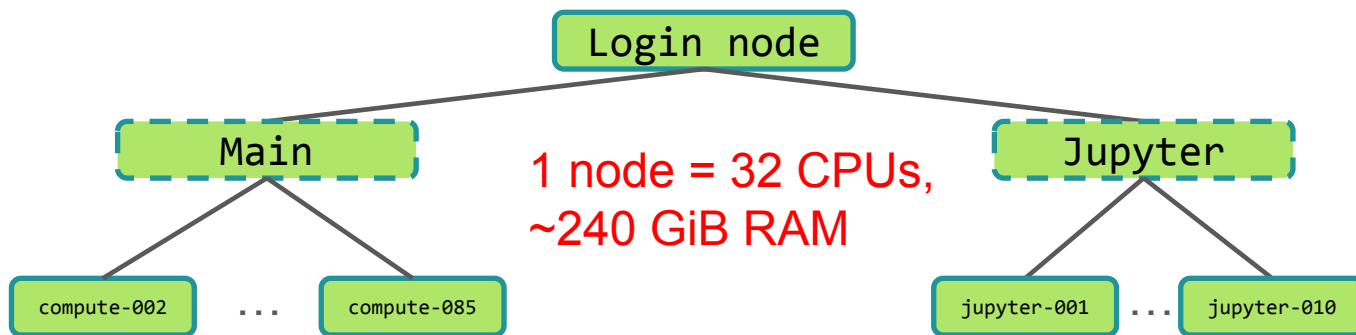
## Session 3: Parallelism

16 September 2025


Dane Kennedy, ilifu Bioinformatics Support  
University of Cape Town



- Login node (job submission & management)
  - Where you land when you log in (also known as “head node”)
  - Run **SLURM commands/submit jobs**, but not software/heavy processes
- Compute nodes
  - Where your processes run (also known as “worker nodes”)
  - Via modules /Singularity containers



- Partitions (other than Jupyter) – see with ‘sinfo’:
  - Main + Devel: 85ish nodes, each w/ 32 CPUs, ~240 GB (usable) RAM
  - HighMem: 2 nodes, w/ 32 CPUs, 503 GB (usable) RAM + 96 CPU 1.5 TB RAM
  - GPU: 7 nodes (P100, V100,..), each w/ 24-48 CPUs, 232-354 GB (usable) RAM



- Oxford definition for parallel processing
  - *a mode of operation in which a process is split into parts, which are executed simultaneously on different processors attached to the same computer [or different computers attached to the same cluster].*
  - A cluster includes many connected nodes, each with its own RAM & CPUs
  - A node = single computer / server / VM / machine / box
- The work is partitioned into smaller jobs, sometimes with a partition of the dataset



# What is a program?

[https://en.wikipedia.org/wiki/Computer\\_program](https://en.wikipedia.org/wiki/Computer_program)

- Set of discrete instructions
- Carried out sequentially
- Example: print average grade of a class

```
1. total = 0
2. for grade in grades:
    total = total + grade
3. average = total / number_of_grades
4. print(average)
```



# Parallel execution of a program

[https://en.wikipedia.org/wiki/Divide-and-conquer\\_algorithm](https://en.wikipedia.org/wiki/Divide-and-conquer_algorithm)

- Partition grades into  $n$  sets and do the following:

```
1. total = 0
2. for grade in 1/n
   grades:

   total = total + grade

1. average_1 = total /
   number_of_grades
```

■ ■ ■

```
1. total = 0
2. for grade in 1/n
   grades:

   total = total + grade

1. average_n = total /
   number_of_grades
```

- Combine results

```
average = (average_1 + ... + average_n)
          number_of_partitions
```

- Executing portions of program simultaneously
- Possible when we have many processors (cores/CPU's)
- Capacity dependent on structure of both hardware AND software
- Requires overall control/coordination mechanism
  - i.e. message passing in MPI / threading / OpenMP

- A cluster includes many connected nodes
- Each node has RAM and multiple cores
- Some nodes have GPUs
- Work of job is partitioned into smaller jobs
- Sometimes with a partition of the data



- Can be achieved on a single machine / node
  - Distributes work over many CPUs
  - Typically implemented using threads / OpenMP
  - GPU
- Or over multiple machines / nodes
  - Distributes work over many tasks, over 1+ nodes
  - Each given amount of memory to use
  - Generally requires a cluster
  - Requires a message passing interface (MPI) wrapper
    - mpirun, srun (SLURM), mpicase (CASA 5)
    - Version of wrapper outside and inside container / venv must match
- Hybrid parallelism? (MPI + OpenMP / MPI + GPU / ...)
- Managed on ilifu by SLURM



- Implementing a normal job in SLURM
  - Will only use 1 CPU, 1 task, and 1 node
  - Default for many processes
- Implementing a threading / OpenMP job in SLURM
  - Need to use >1 CPU, while nodes & tasks must be 1 (unless also using MPI)
    - cpus-per-task (not inherited from #SBATCH)
    - May need to export OMP\_NUM\_THREADS
- Implementing an MPI job in SLURM
  - Need to use >1 task, while nodes and CPUs can be 1
    - nodes, ntasks-per-node, cpus-per-task
    - Best to wrap singularity in MPI call
- Cannot exceed 32 CPUs (or tasks) per node



# SLURM parameters

[https://docs.ilifu.ac.za/#/getting\\_started/submit\\_job\\_slurm?id=customising-your-job-using-sbatchsrn-parameters](https://docs.ilifu.ac.za/#/getting_started/submit_job_slurm?id=customising-your-job-using-sbatchsrn-parameters)

ilifu

<code>--nodes=</code>	<code># the number of nodes allocated to job</code>
<code>--cpus-per-task=</code>	<code># number of cpus per task</code>
<code>--mem-per-cpu=</code>	<code># memory per cpu</code>
<code>--mem=</code>	<code># memory per node</code>
<code>--ntasks-per-node=</code>	<code># number of tasks per node</code>



# SLURM – serial and multi-CPU jobs

- If code is serial, i.e. doesn't use OpenMP or MPI, increasing CPUs or nodes will not decrease execution time

```
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=1

python myscript.py
```

- Using multiple CPUs within a node with OpenMP, where N is an optional number of CPUs (utilised by myscript.py)

```
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=N
#SBATCH --mem-per-cpu=XGB
export OMP_NUM_THREADS=N
python myscript.py
```

- *Note: The maximum number of CPUs per node (32) will not always give the maximum speedup*



- Can also specify tasks or tasks per node

```
#SBATCH --ntasks=N  
#SBATCH --cpus-per-task=1  
#SBATCH --mem=XGB  
  
module add openmpi  
mpirun python myscript.py
```

- Above example doesn't require knowledge of number of node's CPUs; below one does

```
#SBATCH --nodes=1  
#SBATCH --ntasks-per-node=N  
#SBATCH --cpus-per-task=1  
#SBATCH --mem=XGB  
  
module add openmpi  
mpirun python myscript.py
```



- Using multiple nodes with MPI

```
#SBATCH --nodes=N
#SBATCH --ntasks-per-node=n
#SBATCH --cpus-per-task=1
#SBATCH --mem=XGB
module add openmpi
mpirun python myscript.py
```

- *Note: Need to consider that internode communication is slower than intranode communication*
- `--mem` is memory per node, so N times XGB allocated overall (usable by some software)
- Using multiple nodes with MPI as well as multiple cores within node with OpenMP

```
#SBATCH --ntasks=N
#SBATCH --cpus-per-task=n
module add openmpi
mpirun python myscript.py
```



# Live Demo time

[https://github.com/ilifu/ilifu\\_user\\_training/tree/main/session3/tutorial1](https://github.com/ilifu/ilifu_user_training/tree/main/session3/tutorial1)

ilifu



imgflip.com

JAKE-CLARK.TUMBLR

**CBIO**  
Computational Biology @ UCT

**IDiA**



- Don't run software / heavy processes / scp on the login node
  - Only submit jobs and run SLURM commands (sbatch, srun, squeue, etc)
  - Use transfer.ilifu.ac.za to transfer data (external/internal), not login node
- Before running a large job, identify the available resources
  - Use sinfo. Don't hog the cluster. Reduce your allocation if possible
  - Increase likelihood of jobs running with less memory and less walltime
- Use sbatch (srun / screen / tmux / mosh are volatile)
- Cleanup files that aren't needed
  - Old raw data, temporary products, /scratch data, etc
- Don't place large files in your home directory (/users)
- Use Singularity (you cannot install software on the nodes)



# THANK YOU

## Acknowledgements

Dr Jordan Collier for the original slides.



**UNIVERSITY OF CAPE TOWN**  
IYUNIVESITHI YASEKAPA • UNIVERSITEIT VAN KAAPSTAD

